

```

1  """
2  Mask R-CNN
3  The main Mask R-CNN model implemenetation.
4
5  Copyright (c) 2017 Matterport, Inc.
6  Licensed under the MIT License (see LICENSE for details)
7  Written by Waleed Abdulla
8  """
9
10 import os
11 import sys
12 import glob
13 import random
14 import math
15 import datetime
16 import itertools
17 import json
18 import re
19 import logging
20 # from collections import OrderedDict
21 import numpy as np
22 # import scipy.misc
23 import tensorflow as tf
24 # import keras
25 # import keras.backend as K
26 # import keras.layers as KL
27 # import keras.initializers as KI
28 import keras.engine as KE
29 # import keras.models as KM
30 sys.path.append('..')
31 import mrcnn.utils as utils
32
33
34 #####
35 # Detection Target Layer
36 #####
37
38 def overlaps_graph(boxes1, boxes2):
39     """
40     Computes IoU overlaps between two sets of boxes.
41     boxes1, boxes2: [N, (y1, x1, y2, x2)].
42     """
43     # 1. Tile boxes2 and repeate boxes1. This allows us to compare
44     # every boxes1 against every boxes2 without loops.
45     # TF doesn't have an equivalent to np.repeat() so simulate it
46     # using tf.tile() and tf.reshape.
47     b1 = tf.reshape(tf.tile(tf.expand_dims(boxes1, 1),
48                             [1, 1, tf.shape(boxes2)[0]], [-1, 4]))
49     b2 = tf.tile(boxes2, [tf.shape(boxes1)[0], 1])
50     # 2. Compute intersections
51     b1_y1, b1_x1, b1_y2, b1_x2 = tf.split(b1, 4, axis=1)
52     b2_y1, b2_x1, b2_y2, b2_x2 = tf.split(b2, 4, axis=1)
53     y1 = tf.maximum(b1_y1, b2_y1)
54     x1 = tf.maximum(b1_x1, b2_x1)
55     y2 = tf.minimum(b1_y2, b2_y2)
56     x2 = tf.minimum(b1_x2, b2_x2)
57     intersection = tf.maximum(x2 - x1, 0) * tf.maximum(y2 - y1, 0)
58     # 3. Compute unions
59     b1_area = (b1_y2 - b1_y1) * (b1_x2 - b1_x1)
60     b2_area = (b2_y2 - b2_y1) * (b2_x2 - b2_x1)
61     union = b1_area + b2_area - intersection
62     # 4. Compute IoU and reshape to [boxes1, boxes2]
63     iou = intersection / union
64     overlaps = tf.reshape(iou, [tf.shape(boxes1)[0], tf.shape(boxes2)[0]])
65     return overlaps
66
67
68 def detection_targets_graph(proposals, gt_class_ids, gt_boxes, gt_masks, config):
69     """
70     Generates detection targets for one image. Subsamples proposals and
71     generates target class IDs, bounding box deltas, and masks for each.
72
73     Inputs:
74     -----
75     proposals:          [N, (y1, x1, y2, x2)] in normalized coordinates.
76                         Might be zero padded if there are not enough proposals.
77     gt_class_ids:       [MAX_GT_INSTANCES] int class IDs
78     gt_boxes:           [MAX_GT_INSTANCES, (y1, x1, y2, x2)] in normalized coordinates.
79     gt_masks:           [height, width, MAX_GT_INSTANCES] of boolean type.

```

```

80
81     Returns:                Target ROIs and corresponding class IDs, bounding box shifts, and masks.
82     -----
83     rois:                    [TRAIN_ROIS_PER_IMAGE, (y1, x1, y2, x2)] in normalized coordinates
84     class_ids:               [TRAIN_ROIS_PER_IMAGE]. Integer class IDs. Zero padded.
85     deltas:                 [TRAIN_ROIS_PER_IMAGE, NUM_CLASSES, (dy, dx, log(dh), log(dw))]
86                             Class-specific bbox refinements.
87     masks:                  [TRAIN_ROIS_PER_IMAGE, height, width). Masks cropped to bbox
88                             boundaries and resized to neural network output size.
89
90     Note: Returned arrays might be zero padded if not enough target ROIs.
91     """
92     # Assertions
93     asserts = [
94         tf.Assert(tf.greater(tf.shape(proposals)[0], 0), [proposals], name="roi_assertion"),
95     ]
96
97     with tf.control_dependencies(asserts):
98         proposals = tf.identity(proposals)
99
100    # Remove zero padding
101
102    proposals, _ = utils.trim_zeros_graph(proposals, name="trim_proposals")
103    gt_boxes, non_zeros = utils.trim_zeros_graph(gt_boxes, name="trim_gt_boxes")
104    gt_class_ids = tf.boolean_mask(gt_class_ids, non_zeros, name="trim_gt_class_ids")
105    gt_masks = tf.gather(gt_masks, tf.where(non_zeros)[:, 0], axis=2, name="trim_gt_masks")
106
107    # Handle COCO crowds
108    # A crowd box in COCO is a bounding box around several instances. Exclude
109    # them from training. A crowd box is given a negative class ID.
110
111    # tf.where : returns the coordinates of true elements of the specified conditon.
112    #           The coordinates are returned in a 2-D tensor where the first dimension (rows)
113    #           represents the number of true elements, and the second dimension (columns)
114    #           represents the coordinates of the true elements.
115    #           Keep in mind, the shape of the output tensor can vary depending on how many
116    #           true values there are in input. Indices are output in row-major order.
117
118    # tf.gather: Gather slices from params axis (default = 0) according to indices.
119    #           indices must be an integer tensor of any dimension (usually 0-D or 1-D).
120    #           Produces an output tensor with shape params.shape[:axis] + indices.shape + params.shape[axis +
121    1:] where:
122    crowd_ix = tf.where(gt_class_ids < 0)[:, 0]
123    non_crowd_ix = tf.where(gt_class_ids > 0)[:, 0]
124    crowd_boxes = tf.gather(gt_boxes, crowd_ix)
125    crowd_masks = tf.gather(gt_masks, crowd_ix, axis=2)
126    gt_class_ids = tf.gather(gt_class_ids, non_crowd_ix)
127    gt_boxes = tf.gather(gt_boxes, non_crowd_ix)
128    gt_masks = tf.gather(gt_masks, non_crowd_ix, axis=2)
129
130    # Compute overlaps matrix [proposals, gt_boxes]
131    overlaps = overlaps_graph(proposals, gt_boxes)
132
133    # Compute overlaps with crowd boxes [anchors, crowds]
134    crowd_overlaps = overlaps_graph(proposals, crowd_boxes)
135    crowd_iou_max = tf.reduce_max(crowd_overlaps, axis=1)
136    no_crowd_bool = (crowd_iou_max < 0.001)
137
138    # Determine postive and negative ROIs
139    roi_iou_max = tf.reduce_max(overlaps, axis=1)
140
141    # 1. Positive ROIs are those with >= 0.5 IoU with a GT box
142    positive_roi_bool = (roi_iou_max >= 0.5)
143    positive_indices = tf.where(positive_roi_bool)[:, 0]
144
145    # 2. Negative ROIs are those with < 0.5 with every GT box. Skip crowds.
146    negative_indices = tf.where(tf.logical_and(roi_iou_max < 0.5, no_crowd_bool))[:, 0]
147
148    # Subsample ROIs. Aim for 33% positive
149    # Positive ROIs
150    positive_count = int(config.TRAIN_ROIS_PER_IMAGE * config.ROI_POSITIVE_RATIO)
151    positive_indices = tf.random_shuffle(positive_indices)[:positive_count]
152    positive_count = tf.shape(positive_indices)[0]
153
154    # Negative ROIs. Add enough to maintain positive:negative ratio.
155    # negative_count = int((positive_count / config.ROI_POSITIVE_RATIO) - positive_count)
156
157    r = 1.0 / config.ROI_POSITIVE_RATIO

```

```

158     negative_count = tf.cast(r * tf.cast(positive_count, tf.float32), tf.int32) - positive_count
159
160     negative_indices = tf.random_shuffle(negative_indices)[:negative_count]
161
162     # Gather selected ROIs
163     positive_rois = tf.gather(proposals, positive_indices)
164     negative_rois = tf.gather(proposals, negative_indices)
165
166     # Assign positive ROIs to GT boxes.
167     positive_overlaps = tf.gather(overlaps, positive_indices)
168     roi_gt_box_assignment = tf.argmax(positive_overlaps, axis=1)
169     roi_gt_boxes = tf.gather(gt_boxes, roi_gt_box_assignment)
170     roi_gt_class_ids = tf.gather(gt_class_ids, roi_gt_box_assignment)
171
172     # Compute bbox refinement for positive ROIs
173     deltas = utils.box_refinement_graph(positive_rois, roi_gt_boxes)
174     deltas /= config.BBOX_STD_DEV
175
176     # Assign positive ROIs to GT masks
177     # Permute masks to [N, height, width, 1]
178
179     transposed_masks = tf.expand_dims(tf.transpose(gt_masks, [2, 0, 1]), -1)
180
181     # Pick the right mask for each ROI
182     roi_masks = tf.gather(transposed_masks, roi_gt_box_assignment)
183
184     # Compute mask targets
185     boxes = positive_rois
186     if config.USE_MINI_MASK:
187         # Transform ROI corrdinates from normalized image space
188         # to normalized mini-mask space.
189         y1, x1, y2, x2 = tf.split(positive_rois, 4, axis=1)
190         gt_y1, gt_x1, gt_y2, gt_x2 = tf.split(roi_gt_boxes, 4, axis=1)
191         gt_h = gt_y2 - gt_y1
192         gt_w = gt_x2 - gt_x1
193         y1 = (y1 - gt_y1) / gt_h
194         x1 = (x1 - gt_x1) / gt_w
195         y2 = (y2 - gt_y1) / gt_h
196         x2 = (x2 - gt_x1) / gt_w
197         boxes = tf.concat([y1, x1, y2, x2], 1)
198     box_ids = tf.range(0, tf.shape(roi_masks)[0])
199     masks = tf.image.crop_and_resize(tf.cast(roi_masks, tf.float32), boxes,
200                                     box_ids,
201                                     config.MASK_SHAPE)
202     # Remove the extra dimension from masks.
203     masks = tf.squeeze(masks, axis=3)
204
205     # Threshold mask pixels at 0.5 to have GT masks be 0 or 1 to use with
206     # binary cross entropy loss.
207     masks = tf.round(masks)
208
209     # Append negative ROIs and pad bbox deltas and masks that
210     # are not used for negative ROIs with zeros.
211     rois = tf.concat([positive_rois, negative_rois], axis=0)
212     N = tf.shape(negative_rois)[0]
213     P = tf.maximum(config.TRAIN_ROIS_PER_IMAGE - tf.shape(rois)[0], 0)
214     rois = tf.pad(rois, [(0, P), (0, 0)])
215     roi_gt_boxes = tf.pad(roi_gt_boxes, [(0, N + P), (0, 0)])
216     roi_gt_class_ids = tf.pad(roi_gt_class_ids, [(0, N + P)])
217     deltas = tf.pad(deltas, [(0, N + P), (0, 0)])
218     masks = tf.pad(masks, [[0, N + P], (0, 0), (0, 0)])
219
220     return rois, roi_gt_class_ids, deltas, masks
221
222
223 class DetectionTargetLayer(KE.Layer):
224     """Subsamples proposals and generates target box refinement, class_ids,
225     and masks for each.
226
227     Inputs:
228     -----
229     proposals: [batch, N, (y1, x1, y2, x2)] in normalized coordinates. Might
230                be zero padded if there are not enough proposals.
231     gt_class_ids: [batch, MAX_GT_INSTANCES] Integer class IDs.
232     gt_boxes: [batch, MAX_GT_INSTANCES, (y1, x1, y2, x2)] in normalized
233               coordinates.
234     gt_masks: [batch, height, width, MAX_GT_INSTANCES] of boolean type
235
236     Returns:

```

```

237         -----
238         Target ROIs and corresponding class IDs, bounding box shifts, and masks.
239         rois: [batch, TRAIN_ROIS_PER_IMAGE, (y1, x1, y2, x2)] in normalized coordinates
240         target_class_ids: [batch, TRAIN_ROIS_PER_IMAGE]. Integer class IDs.
241         target_deltas: [batch, TRAIN_ROIS_PER_IMAGE, NUM_CLASSES, (dy, dx, log(dh), log(dw), class_id)]
242                        Class-specific bbox refinements.
243         target_mask: [batch, TRAIN_ROIS_PER_IMAGE, height, width)
244                      Masks cropped to bbox boundaries and resized to neural network output size.
245
246         Note: Returned arrays might be zero padded if not enough target ROIs.
247         """
248
249         def __init__(self, config, **kwargs):
250             # super(DetectionTargetLayer, self).__init__(**kwargs)
251             super().__init__(**kwargs)
252             self.config = config
253
254         def call(self, inputs):
255             proposals = inputs[0]
256             gt_class_ids = inputs[1]
257             gt_boxes = inputs[2]
258             gt_masks = inputs[3]
259
260             # Slice the batch and run a graph for each slice
261             # TODO: Rename target_bbox to target_deltas for clarity
262
263             names = ["rois", "target_class_ids", "target_bbox", "target_mask"]
264             outputs = utils.batch_slice([proposals, gt_class_ids, gt_boxes, gt_masks],
265                                         lambda w, x, y, z: detection_targets_graph(w, x, y, z, self.config),
266                                         self.config.IMAGES_PER_GPU, names=names)
267             return outputs
268
269         def compute_output_shape(self, input_shape):
270             return [
271                 (None, self.config.TRAIN_ROIS_PER_IMAGE, 4), # rois
272                 (None, 1), # class_ids
273                 (None, self.config.TRAIN_ROIS_PER_IMAGE, 4), # deltas
274                 (None, self.config.TRAIN_ROIS_PER_IMAGE, self.config.MASK_SHAPE[0],
275                  self.config.MASK_SHAPE[1]) # masks
276             ]
277
278         def compute_mask(self, inputs, mask=None):
279             return [None, None, None, None]
280
281
282         #####
283         # Detection Layer
284         #####
285
286         def clip_to_window(window, boxes):
287             """
288             window: (y1, x1, y2, x2). The window in the image we want to clip to.
289             boxes: [N, (y1, x1, y2, x2)]
290             """
291             boxes[:, 0] = np.maximum(np.minimum(boxes[:, 0], window[2]), window[0])
292             boxes[:, 1] = np.maximum(np.minimum(boxes[:, 1], window[3]), window[1])
293             boxes[:, 2] = np.maximum(np.minimum(boxes[:, 2], window[2]), window[0])
294             boxes[:, 3] = np.maximum(np.minimum(boxes[:, 3], window[3]), window[1])
295             return boxes
296
297
298         def refine_detections(rois, probs, deltas, window, config):
299             """Refine classified proposals and filter overlaps and return final
300             detections.
301
302             Inputs:
303             rois: [N, (y1, x1, y2, x2)] in normalized coordinates
304             probs: [N, num_classes]. Class probabilities.
305             deltas: [N, num_classes, (dy, dx, log(dh), log(dw))]. Class-specific
306                    bounding box deltas.
307             window: (y1, x1, y2, x2) in image coordinates. The part of the image
308                    that contains the image excluding the padding.
309
310             Returns detections shaped: [N, (y1, x1, y2, x2, class_id, score)]
311             """
312             # Class IDs per ROI
313             class_ids = np.argmax(probs, axis=1)
314             # Class probability of the top class of each ROI
315             class_scores = probs[np.arange(class_ids.shape[0]), class_ids]

```

```

316     # Class-specific bounding box deltas
317     deltas_specific = deltas[np.arange(deltas.shape[0]), class_ids]
318     # Apply bounding box deltas
319     # Shape: [boxes, (y1, x1, y2, x2)] in normalized coordinates
320     refined_rois = utils.apply_box_deltas(
321         rois, deltas_specific * config.BBOX_STD_DEV)
322     # Convert coordinates to image domain
323     # TODO: better to keep them normalized until later
324     height, width = config.IMAGE_SHAPE[:2]
325     refined_rois *= np.array([height, width, height, width])
326     # Clip boxes to image window
327     refined_rois = clip_to_window(window, refined_rois)
328     # Round and cast to int since we're dealing with pixels now
329     refined_rois = np.rint(refined_rois).astype(np.int32)
330
331     # TODO: Filter out boxes with zero area
332
333     # Filter out background boxes
334     keep = np.where(class_ids > 0)[0]
335     # Filter out low confidence boxes
336     if config.DETECTION_MIN_CONFIDENCE:
337         keep = np.intersect1d(
338             keep, np.where(class_scores >= config.DETECTION_MIN_CONFIDENCE)[0])
339
340     # Apply per-class NMS
341     pre_nms_class_ids = class_ids[keep]
342     pre_nms_scores = class_scores[keep]
343     pre_nms_rois = refined_rois[keep]
344     nms_keep = []
345     for class_id in np.unique(pre_nms_class_ids):
346         # Pick detections of this class
347         ix = np.where(pre_nms_class_ids == class_id)[0]
348         # Apply NMS
349         class_keep = utils.non_max_suppression(
350             pre_nms_rois[ix], pre_nms_scores[ix],
351             config.DETECTION_NMS_THRESHOLD)
352         # Map indices
353         class_keep = keep[ix[class_keep]]
354         nms_keep = np.union1d(nms_keep, class_keep)
355     keep = np.intersect1d(keep, nms_keep).astype(np.int32)
356
357     # Keep top detections
358     roi_count = config.DETECTION_MAX_INSTANCES
359     top_ids = np.argsort(class_scores[keep])[::-1][:roi_count]
360     keep = keep[top_ids]
361
362     # Arrange output as [N, (y1, x1, y2, x2, class_id, score)]
363     # Coordinates are in image domain.
364     result = np.hstack((refined_rois[keep],
365                         class_ids[keep][..., np.newaxis],
366                         class_scores[keep][..., np.newaxis]))
367     return result
368
369
370 class DetectionLayer(KE.Layer):
371     """Takes classified proposal boxes and their bounding box deltas and
372     returns the final detection boxes.
373
374     Returns:
375     [batch, num_detections, (y1, x1, y2, x2, class_score)] in pixels
376     """
377
378     def __init__(self, config=None, **kwargs):
379         super(DetectionLayer, self).__init__(**kwargs)
380         self.config = config
381
382     def call(self, inputs):
383         def wrapper(rois, mrcnn_class, mrcnn_bbox, image_meta):
384             detections_batch = []
385             for b in range(self.config.BATCH_SIZE):
386                 _, _, window, _ = parse_image_meta(image_meta)
387                 detections = refine_detections(
388                     rois[b], mrcnn_class[b], mrcnn_bbox[b], window[b], self.config)
389                 # Pad with zeros if detections < DETECTION_MAX_INSTANCES
390                 gap = self.config.DETECTION_MAX_INSTANCES - detections.shape[0]
391                 assert gap >= 0
392                 if gap > 0:
393                     detections = np.pad(
394                         detections, [(0, gap), (0, 0)], 'constant', constant_values=0)

```

```
395         detections_batch.append(detections)
396
397         # Stack detections and cast to float32
398         # TODO: track where float64 is introduced
399         detections_batch = np.array(detections_batch).astype(np.float32)
400         # Reshape output
401         # [batch, num_detections, (y1, x1, y2, x2, class_score)] in pixels
402         return np.reshape(detections_batch, [self.config.BATCH_SIZE, self.config.DETECTION_MAX_INSTANCES, 6])
403
404     # Return wrapped function
405     return tf.py_func(wrapper, inputs, tf.float32)
406
407 def compute_output_shape(self, input_shape):
408     return (None, self.config.DETECTION_MAX_INSTANCES, 6)
409
410
```